

Donald Trump Tweet Author Analysis

This personal project looks at data from president Trump's Twitter. Donald Trump used to tweet from an iPhone while his campaign staff tweeted through his account on Android, so this data can be used to train a neural network to predict who wrote his newer tweets based on a set of features to include how often he used certain words and how long the tweets were.

Load Data

```
% Uncomment this on first run to load in data, takes a while
% opts = detectImportOptions('dataset.xlsx');
% opts.SelectedVariableNames = [8 12:334];
% X = readmatrix('dataset.xlsx',opts);
% opts.SelectedVariableNames = 2;
% Y = readmatrix('dataset.xlsx',opts);
load('X.mat');
load('Y.mat');
s = string(Y);
id = s == "Twitter for iPhone" | s == "Twitter for Android";
s = s(id);
X = X(id,:);
Y = s == "Twitter for iPhone";
```

Mean normalize X

```
non_present = max(X) == 0;
X = X(:,~non_present);
X = (X - repmat(mean(X), size(X,1), 1)) ./ repmat(2*std(X), size(X,1), 1);
X = [ones(size(X,1),1) X];
```

Divide data

```
m = size(X,1);
[trainInd,valInd,testInd] = dividerand(m,.6,.2,.2);
X_train = X(trainInd,:);
Y_train = Y(trainInd,:);
X_val = X(valInd,:);
Y_val = Y(valInd,:);
X_test = X(testInd,:);
Y_test = Y(testInd,:);
```

Logistic Regression - Train theta

Here I will use logistic regression, then I will try a neural network.

```
initial_theta = zeros(size(X_train,2),1);
options = optimoptions(@fminunc,'Algorithm','Quasi-Newton','GradObj','on','MaxIter',400);
[theta, cost] = fminunc(@(t)(costFunctionReg(t, X_train, Y_train,1)), initial_theta, options);
```

Local minimum found.

Optimization completed because the size of the gradient is less than the value of the optimality tolerance.

<stopping criteria details>

```
p = predict(theta, X_val);
success = p == Y_val;
success = success(success);
success_rate = length(success)/length(p);
fprintf('Logistic Regression Validation Set Accuracy = %f\n', success_rate);
```

Logistic Regression Validation Set Accuracy = 0.850289

Neural Network

The neural network is a simple 1-hidden layer neural network.

Initialize Thetas

```
% nn script
hidden_layer_size = 40;
input_layer_size = 308;

initial_Theta1 = randInitializeWeights(input_layer_size, hidden_layer_size);
initial_Theta2 = randInitializeWeights(hidden_layer_size, 1);

% Unroll parameters
initial_nn_params = [initial_Theta1(:) ; initial_Theta2(:)];
```

Train

```
options = optimset('MaxIter', 60);
lambda = 3;

% Create "short hand" for the cost function to be minimized
costFunction = @(p) nnCostFunction(p, input_layer_size, hidden_layer_size, X_train, Y_train, lambda);

% Now, costFunction is a function that takes in only one argument (the
% neural network parameters)
[nn_params, ~] = fmincg(costFunction, initial_nn_params, options);
```

```
Iteration    1 | Cost: 6.980162e-01
Iteration    2 | Cost: 4.951659e-01
Iteration    3 | Cost: 4.478344e-01
Iteration    4 | Cost: 4.135613e-01
Iteration    5 | Cost: 3.823161e-01
Iteration    6 | Cost: 3.513519e-01
Iteration    7 | Cost: 3.427668e-01
Iteration    8 | Cost: 3.306680e-01
Iteration    9 | Cost: 3.296712e-01
Iteration   10 | Cost: 3.237273e-01
Iteration   11 | Cost: 3.203762e-01
Iteration   12 | Cost: 3.155860e-01
Iteration   13 | Cost: 3.138944e-01
Iteration   14 | Cost: 3.133712e-01
```

| | | |
|-----------|----|--------------------|
| Iteration | 15 | Cost: 3.102120e-01 |
| Iteration | 16 | Cost: 3.082352e-01 |
| Iteration | 17 | Cost: 3.047811e-01 |
| Iteration | 18 | Cost: 3.017493e-01 |
| Iteration | 19 | Cost: 2.973636e-01 |
| Iteration | 20 | Cost: 2.968866e-01 |
| Iteration | 21 | Cost: 2.952505e-01 |
| Iteration | 22 | Cost: 2.940775e-01 |
| Iteration | 23 | Cost: 2.935262e-01 |
| Iteration | 24 | Cost: 2.929661e-01 |
| Iteration | 25 | Cost: 2.921174e-01 |
| Iteration | 26 | Cost: 2.920055e-01 |
| Iteration | 27 | Cost: 2.893992e-01 |
| Iteration | 28 | Cost: 2.861177e-01 |
| Iteration | 29 | Cost: 2.852148e-01 |
| Iteration | 30 | Cost: 2.846142e-01 |
| Iteration | 31 | Cost: 2.845092e-01 |
| Iteration | 32 | Cost: 2.842318e-01 |
| Iteration | 33 | Cost: 2.838166e-01 |
| Iteration | 34 | Cost: 2.836979e-01 |
| Iteration | 35 | Cost: 2.824089e-01 |
| Iteration | 36 | Cost: 2.809137e-01 |
| Iteration | 37 | Cost: 2.763423e-01 |
| Iteration | 38 | Cost: 2.739081e-01 |
| Iteration | 39 | Cost: 2.737138e-01 |
| Iteration | 40 | Cost: 2.694333e-01 |
| Iteration | 41 | Cost: 2.581496e-01 |
| Iteration | 42 | Cost: 2.551452e-01 |
| Iteration | 43 | Cost: 2.516281e-01 |
| Iteration | 44 | Cost: 2.514009e-01 |
| Iteration | 45 | Cost: 2.510815e-01 |
| Iteration | 46 | Cost: 2.497073e-01 |
| Iteration | 47 | Cost: 2.491823e-01 |
| Iteration | 48 | Cost: 2.471236e-01 |
| Iteration | 49 | Cost: 2.463393e-01 |
| Iteration | 50 | Cost: 2.451785e-01 |
| Iteration | 51 | Cost: 2.450664e-01 |
| Iteration | 52 | Cost: 2.436585e-01 |
| Iteration | 53 | Cost: 2.429701e-01 |
| Iteration | 54 | Cost: 2.426089e-01 |
| Iteration | 55 | Cost: 2.402617e-01 |
| Iteration | 56 | Cost: 2.346075e-01 |
| Iteration | 57 | Cost: 2.329964e-01 |
| Iteration | 58 | Cost: 2.254908e-01 |
| Iteration | 59 | Cost: 2.252035e-01 |
| Iteration | 60 | Cost: 2.240138e-01 |

```

Theta1 = reshape(nn_params(1:hidden_layer_size * (input_layer_size + 1)), hidden_layer_size, (input_layer_size + 1));
Theta2 = reshape(nn_params((1 + (hidden_layer_size * (input_layer_size + 1))):end), 1, (hidden_layer_size + 1));

pred_val = predict_nn(Theta1, Theta2, X_val);
val_accuracy = mean(double(pred_val == Y_val));
fprintf('Logistic Regression Validation Set Accuracy = %f\n', val_accuracy);

```

Logistic Regression Validation Set Accuracy = 0.865171